

# Automatically Stopping Idle Instance Using Lambda and CloudWatch

## Save AWS EC2 Cost by Automatically Stopping Idle Instance Using Lambda and CloudWatch

In machine learning, especially deep learning, GPU is often used to train machine learning models. In most cases, GPU instances on the cloud services are selected by engineers because they are easier to prepare. But since the GPU instance is expensive, it is a waste of cost if you leave it running when the state of the instance is idle. In such a case, you can significantly reduce the cost by automatically stopping the instance.

First, we will attach a CloudWatch alarm to an EC2 instance when the state of the instance becomes "running". The alarm monitors some metrics and automatically stops the instance when the alarm is raised. In order to implement these behaviors, we will use AWS Lambda and CloudWatch. Both services are serverless, so we don't have to care about their operation and maintenance. The overview is as follows:



This tutorial is divided into 3 parts; they are:

- Create an IAM role
- Create Lambda function
- Create CloudWatch Event

Let's get started.

### Create an IAM Role

Let's create a role that can access to EC2 instance from the Lambda function.

Move to the "Create IAM role" page and select Lambda in "Choose the service that will use this role" and click "Next: Permissions" button:

#### Choose the service that will use this role

##### EC2

Allows EC2 instances to call AWS services on your behalf.

##### Lambda

Allows Lambda functions to call AWS services on your behalf.

API Gateway	Config	EMR	IoT	Rekognition
AWS Support	DMS	ElastiCache	Kinesis	S3
AppSync	Data Lifecycle Manager	Elastic Beanstalk	Lambda	SMS
Application Auto Scaling	Data Pipeline	Elastic Container Service	Lex	SNS
Auto Scaling	DeepLens	Elastic Transcoder	Machine Learning	SWF
Batch	Directory Service	ElasticLoadBalancing	Macie	SageMaker
CloudFormation	DynamoDB	Glue	MediaConvert	Service Catalog
CloudHSM	EC2	Greengrass	OpsWorks	Step Functions
CloudWatch Events	EC2 - Fleet	GuardDuty	RDS	Storage Gateway
CodeBuild	EKS	Inspector	Redshift	Trusted Advisor
CodeDeploy				

\* Required









Cancel

Next: Permissions

Then select the policy to attach to the role. Since I want to access to the EC2 instance from Lambda function, I will attach "Amazon EC2 Full Access" policy. After selecting the policy, click "Next: Review" button. If you want to set only the minimum permission, Create a policy yourself and attach it

Filter policies ▾

Q EC2

		Policy name ▾	Used as
<input type="checkbox"/>	▶	 AmazonEC2ContainerServiceforEC2Role	None
<input type="checkbox"/>	▶	 AmazonEC2ContainerServiceFullAccess	None
<input type="checkbox"/>	▶	 AmazonEC2ContainerServiceRole	Permissions policy (1)
<input checked="" type="checkbox"/>	▶	 AmazonEC2FullAccess	Permissions policy (1)
<input type="checkbox"/>	▶	 AmazonEC2ReadOnlyAccess	None
<input type="checkbox"/>	▶	 AmazonEC2ReportsAccess	None
<input type="checkbox"/>	▶	 AmazonEC2RoleforAWSCodeDeploy	None
<input type="checkbox"/>	▶	 AmazonEC2RoleforDataPipelineRole	None

After input the role name, select "Create role" button to create the role. Here, the role is named "AWSLambdaEC2FullAccess".

✓ The role **AWSEC2CostReduction** has been created.

### Create a Lambda function

Next we will create a Lambda function. In the Lambda function, we do two things:

- Get the EC2 instance ID
- Create an alarm and attach it to the EC2 instance

After moving to the Lambda page, you should set the function setting as follows. Here, you must select the previously created role as existing role:

**Author from scratch** [Info](#)

Name

Runtime

Role

Defines the permissions of your function. Note that new roles may not be available for a few minutes after creation. [Learn more](#) execution roles.

Existing role

You may use an existing role with this function. Note that the role must be assumable by Lambda and must have Cloudwatch Logs permissions.

After creating the function, we need to write some code in Python. Use boto3 package to get the instance id and create the CloudWatch alarm and attach it to the instance. It can be written as follows.

```
import boto3

def put_cpu_alarm(instance_id):
    cloudWatch = boto3.client('cloudwatch')
    cloudWatch.put_metric_alarm(
        AlarmName = f'CPU_ALARM_{instance_id}',
        AlarmDescription = 'Alarm when server CPU does not exceed 10%',
        AlarmActions = ['arn:aws:automate:us-east-1:ec2:stop'],
        MetricName = 'CPUUtilization',
        Namespace = 'AWS/EC2',
        Statistic = 'Average',
        Dimensions = [{'Name': 'InstanceId', 'Value': instance_id}],
        Period = 300,
        EvaluationPeriods = 3,
        Threshold = 10,
        ComparisonOperator = 'LessThanOrEqualToThreshold',
        TreatMissingData = 'notBreaching'
    )

def lambda_handler(event, context):
    instance_id = event['detail']['instance-id']
    ec2 = boto3.resource('ec2')
    instance = ec2.Instance(instance_id)
    if instance.instance_type.endswith('xlarge'):
        put_cpu_alarm(instance_id)
```

put\_cpu\_alarm function creates a CloudWatch alarm and attach it to the specified instance. The alarm monitors the CPU utilization on the instance and stops the instance when it reaches the utilization threshold (below 10%) 3 times in a row. Let's change this setting according to your own environment. I think that the materials listed below are useful.

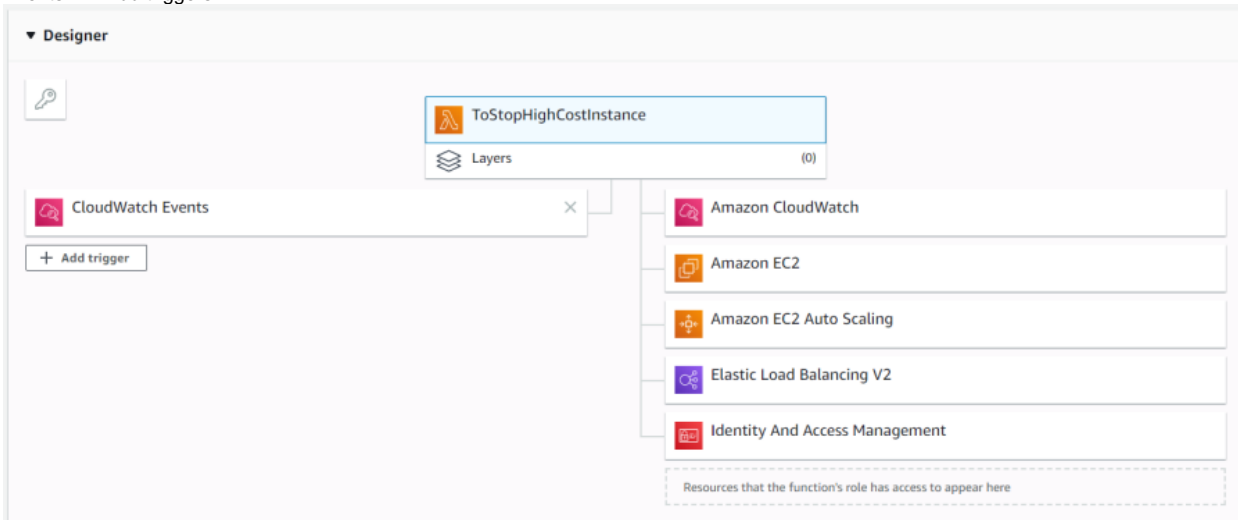
- [Create an Amazon CloudWatch alarm](#)
- [Creating Alarms in Amazon CloudWatch](#)
- [Docs/Available Services/CloudWatch](#)

**lambda\_handler** is the entry point to be invoked when it detects the running of the instance. Internally, we get the instance type, and call **put\_cpu\_alarm** if the instance type name is ends with xlarge. If you want to attach the alarm only to a specific instance type, change the condition in if statement. In this time, I selected CPU utilization as the monitoring metrics, but you can also monitor GPU utilization by adding custom metrics. The method is detailed in the following documents.

- [Monitoring GPU utilization with Amazon CloudWatch](#)

### Create CloudWatch Events

Finally, let's set up to call the Lambda function when we detect running the instance. After moving to the Lambda function page, let's select "CloudWatch Events" in "Add triggers".



After adding the trigger, select the CloudWatch Events you added. Then you will see a form for setting the trigger at the bottom of the page. After selecting "Create New Rule", set it as follows. By setting like this, you can call the Lambda function when the instance starts up.

#### Rule type

Trigger your target based on an event pattern, or based on an automated schedule.

- ☒ Event pattern
- ☐ Schedule expression

EC2

EC2 instance state-change notification

☒ Detail

#### Detail

☒ State ☐ Instances

#### State

× running

After executing the lambda function we will get the output as shown in below.

ToStopHighCostInstance
Throttle
Qualifiers
Actions
Test
Test
Save

Execution result: succeeded (logs)

Details

The area below shows the result returned by your function execution. [Learn more](#) about returning results from your function.

null

Summary

Code SHA-256  
IdRyocKmyDpd+Ch3vMKkfq5gpsrvImnwnJP3YjrsoqM=

Request ID  
787cd27f-416d-4523-8dc6-b6cbafcbcdcb

Duration  
2466.39 ms

Billed duration  
2500 ms

Resources configured  
128 MB

Max memory used  
82 MB Init Duration: 185.29 ms

Log output

The section below shows the logging calls in your code. These correspond to a single row within the CloudWatch log group corresponding to this Lambda function. [Click here](#) to view the CloudWatch log group.

```

START RequestId: 787cd27f-416d-4523-8dc6-b6cbafcbcdcb Version: $LATEST
END RequestId: 787cd27f-416d-4523-8dc6-b6cbafcbcdcb
REPORT RequestId: 787cd27f-416d-4523-8dc6-b6cbafcbcdcb Duration: 2466.39 ms Billed Duration: 2500 ms Memory Size: 128 MB Max Memory Used: 82 MB Init Duration: 185.29 ms

```

The below output shows the history of EC2 machine alarms how those changed from critical stage to OK state.

History (9)

< 1 2 >

Date	Type	Description
2019-11-06 11:06:35	State update	Alarm updated from <b>In alarm</b> to <b>OK</b>
2019-11-06 10:55:35	Action	Successfully executed action <code>arn:aws:automate:us-east-1:ec2:stop</code>
2019-11-06 10:55:35	Action	Successfully executed action <code>arn:aws:automate:us-east-1:ec2:stop</code>
2019-11-06 10:55:35	State update	Alarm updated from <b>OK</b> to <b>In alarm</b>
2019-11-06 09:05:35	State update	Alarm updated from <b>In alarm</b> to <b>OK</b>
2019-11-06 08:57:36	Action	Successfully executed action <code>arn:aws:automate:us-east-1:ec2:stop</code>
2019-11-06 08:57:35	Action	Successfully executed action <code>arn:aws:automate:us-east-1:ec2:stop</code>
2019-11-06 08:57:35	State update	Alarm updated from <b>Insufficient data</b> to <b>In alarm</b>

Once the lambda triggered after the given time intervals, EC2 will be stopped.

i-0...5
t2.micro
us-east-1a
stopped
OK

## Summary

In this tutorial, you discovered how to automatically stop an EC2 instance when the state is idle. Specifically, you learned:

- How to create an IAM role.
- How to create a Lambda function.
- How to create CloudWatch alarm.

Do you have any questions?

Ask your questions in the comments below and I will do my best to answer.

## Reference

- [Create an Amazon CloudWatch alarm](#)
- [Creating Alarms in Amazon CloudWatch](#)

- [Docs/Available Services/CloudWatch](#)
- [Monitoring GPU utilization with Amazon CloudWatch](#)